# CHERI-picking: Leveraging capability hardware for pre-fetching

Shaurya Patel, Sidhartha Agrawal, Alexandra Fedorova, Margo Seltzer

## Problem

- DRAM is **expensive** and **memory demands** of datacenter applications **are increasing**.
- Far memory **is cheap** but **degrades application performance**, prefetchers are important to hide latency of memory access.
- Applications with **Linked data structures** are difficult to prefetch because of their pointer-chasing behavior.

## Insightful data from hardware

- Capability hardware systems such as CHERI makes it possible to identify pointers at runtime.

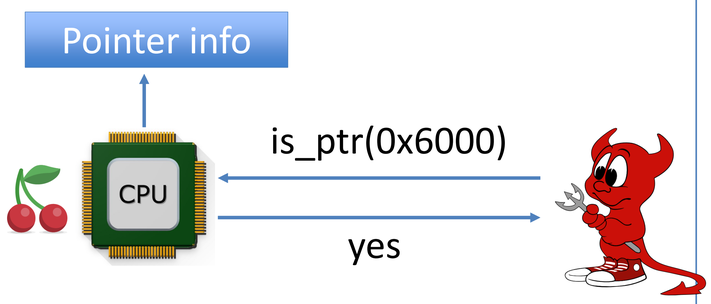


## Our Approach

- Analyze application traces to classify percentage of pagefaults caused by pointer accesses.
- Leverage CHERI data at runtime to make OS pre-fetching decisions.

## Prefetching system for pointer data

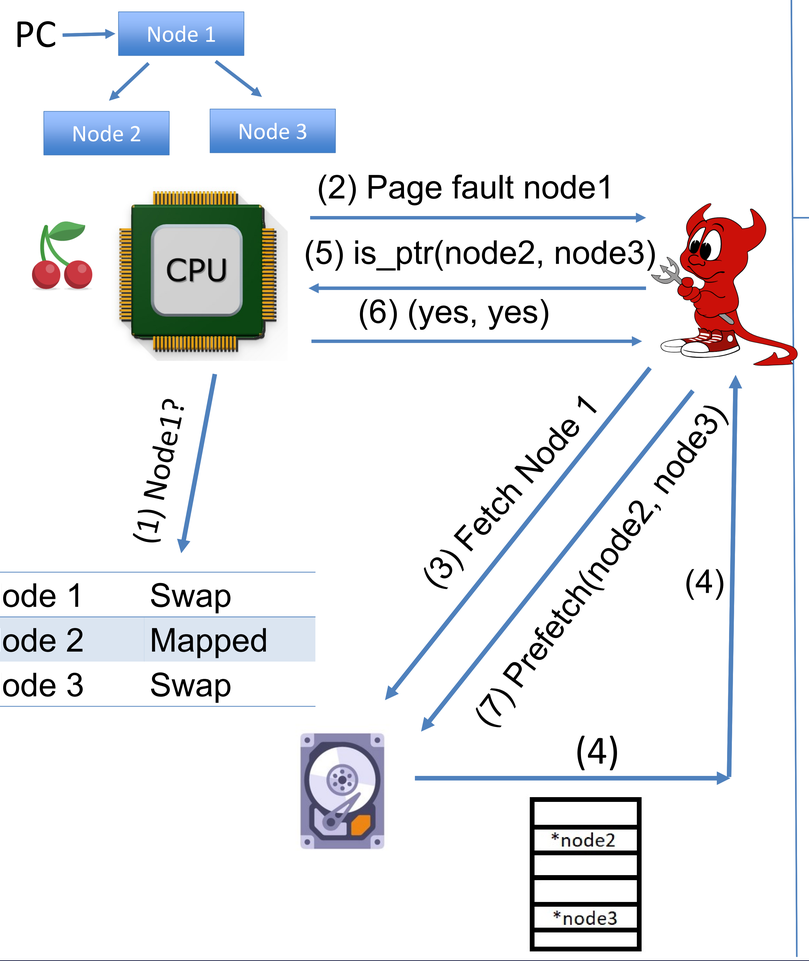


| | |
|---|---|
| Node 1 | Swap |
| Node 2 | Mapped |
| Node 3 | Swap |

## Preliminary results

- We benchmark the system by using a linked list traversal.
- We used a random permutation of numbers to assign linked list nodes, each node is page size.
- Total memory consumption is 4GB and we constrained it at 50%.

| | Default prefetcher | Pointer based prefetcher | Improvement |
|---|---|---|---|
| Cache hits | 46749 | 655345 | 12.9X ↑ |
| #major faults | 1.86m | 1.41m | 24% ↓ |
| Execution time | 21.85 minutes | 20.56 minutes | 5.91% ↑ |

## Future work

- Use application execution traces to identify programs for which pointer-based prefetching can be effective.
- Evaluate the pre-fetcher CHERI-Morello hardware using CheriBSD.
- Develop real-time monitoring to identify programs for which we should do pointer-based prefetching.